

4400 review - CUDA

CUDA hierarchy:

- 1.) threads
 - 2.) blocks
 - 3.) grids
- } Can be 1D, 2D, 3D

CPU: host GPU: device

Kernel is specified with a
--global-- specifier

2.1.2.1. Specifying Kernels

The code for a kernel is specified using the `__global__` declaration specifier. This indicates to the compiler that this function will be compiled for the GPU in a way that allows it to be invoked from a kernel launch. A kernel launch is an operation which starts a kernel running, usually from the CPU. Kernels are functions with a `void` return type.

```
// Kernel definition
__global__ void vecAdd(float* A, float* B, float* C)
{}
```

`vecAdd <<< numBlocks, numThreads>>> (A, B, C)`

will launch `vecAdd` with that many blocks and threads.

Global index of a thread:

1D: $idx = blockDim.x * blockIdx.x + threadIdx.x$

2D: $idx_x = blockDim.x * blockIdx.x + threadIdx.x$

$idx_y = blockDim.y * blockIdx.y + threadIdx.y$

Then use row major ordering to calculate a 1D index on flattened 2D array.

$width = blockDim.x * GridDim.x$

dim3 is a CUDA datatype:

dim3 gridDim	size of grid
dim3 blockDim	size of block
dim3 blockIdx	block index within the grid
dim3 threadIdx	thread index within the block

`cudaDeviceSynchronize()` after each iteration

TECHNICALLY, this is inefficient. But this is the way that was taught to us.